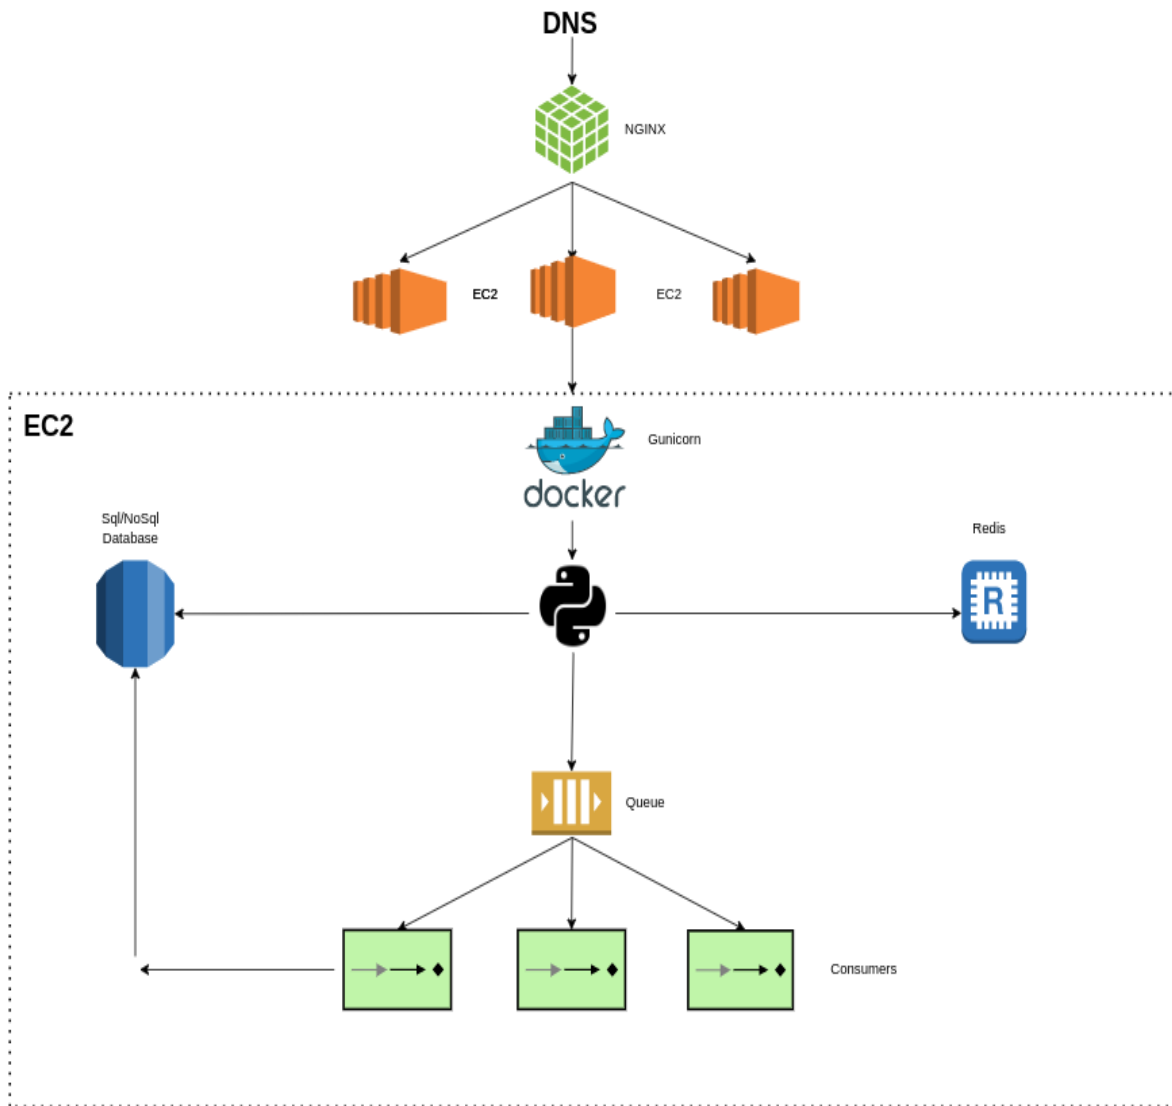


Architecture For Scalable Movie Database



Components:

- NGINX
 - Using Nginx as an Web Server for our application in order to handle the following:
 - Concurrent Connections
 - Static Files (Current Application may include Image, Audio, Video related to Movies)
 - Load Balancer : Efficient redirection of traffic onto lower level Ec2 instances.
 - No Downtime during reload in case of route change.
- Ec2 Instance
 - Ec2 instance to host and run the application itself.
- Docker
 - Docker to run, deploy applications with ease.
- Gunicorn
 - Gunicorn to be used as standard WSGI to forward requests to the application.
 - We can set up multiple worker/thread combinations to distribute incoming requests.
- Python Application
 - Current Application is written in Flask.
 - Although flask < 2.0 does not support async calls, we can use celery with flask to publish long running tasks onto the queue in an async manner.
 - In our application, we have decoupled the `create/update/delete` operations to celery.

- Sql Database
 - A relational database is preferred since we know that the entities to be stored are related to each other.(eg. Actors, Movies, Shows, Genres etc).
 - Scaling can be done Vertically in order to handle increasing load.
 - Schemas are consistent and hence structure is defined well in advance.
- Redis
 - User and token level details can be stored onto Redis for authentication/authorization purposes.
- RabbitMq
 - RabbitMq is used as a Message Broker for queuing messages and their eventual consumption.
- Celery Worker
 - Multiple Celery workers can be spawned for asynchronous publishing and consumption of messages.